

Lane Keep Assist (LKA) System – Final Project Report

Classical Computer Vision Implementation

Author: Kipchirchir Raphael ([LGL7CS](#))

Date: November 2025

1. Project Overview

This project implements a **real-time Lane Keep Assist (LKA)** system using **classical computer vision** techniques. The system processes driving videos to:

- Detect **left and right lane boundaries**
- Estimate **lateral offset** (in meters)
- Generate **annotated video** with colored lane overlays and HUD
- Export **per-frame metrics** to CSV
- Ensure **temporal stability** and **confidence-aware detection**

Implementation Path: Classical CV Baseline (recommended)

No deep learning used – fully interpretable and real-time capable.

2. System Architecture

```
src/
├── main.py           → Pipeline orchestrator
├── preprocess.py    → Color + edge thresholding
├── warp.py          → PerspectiveTransformer (IPM)
├── lane_fit.py       → LaneDetector (sliding window + polyfit)
├── temporal.py      → TemporalFilter (smoothing)
├── overlay.py        → OverlayRenderer (HUD + drawing)
├── metrics.py        → MetricsCalculator (offset + CSV)
└── utils.py          → ROI masking
```

Output:

```
outputs/
├── annotated_output.mp4
└── per_frame_metrics.csv
```

3. Pipeline Step-by-Step

3.1 Input Frame

- Read from video (`cv2.VideoCapture`)
- Full resolution preserved

3.2 Preprocessing (`LanePreprocessor`)

- **Color Spaces:** HLS, LAB, HSV
- **White Lane Detection:**
 - $\text{HLS}[\text{L}] \geq 220$, $\text{HSV}[\text{S}] \leq 30$, $\text{LAB}[\text{L}] \geq 220$
 - Conservative AND logic
- **Edge Detection:**
 - Sobel X on grayscale ($|\nabla x| > 50$)
- **Morphology:**
 - Open (3×3) → Close (5×5)
- **ROI Mask** (trapezoidal, **not cropped**)

3.3 Inverse Perspective Mapping (`PerspectiveTransformer`)

- **Source:** Trapezoid in camera view
 - [0.20w, 0.95h], [0.40w, 0.65h], [0.60w, 0.65h], [0.80w, 0.95h]
- **Destination:** Rectangle in bird's-eye
- **Dynamic Adjustment:** Updates source points if $\text{confidence} > 0.7$

3.4 Lane Pixel Extraction (`LaneDetector`)

- **Histogram:** Bottom half → left/right peaks
- **Sliding Windows:**
 - 9 windows, 100px margin, 50px min pixels
 - Recenters on mean pixel position
- **Output:** Pixel indices + **debug visualization**

3.5 Curve Fitting

- **Polynomial:** `(x = a*y2 + b*y + c)` (`np.polyfit`)
- **Minimum 6 pixels required**
- **Rejects extreme curvature**

3.6 Confidence Scoring

```
conf = 0.5 * (pixel_count / 2000) + 0.5 * (1 - residual / 100)
```

- **Threshold:** `τ = 0.6`
- **Detection Flag:** `conf > τ → YES`

3.7 Temporal Smoothing (`TemporalFilter`)

- **Exponential smoothing:**

```
α = min(confidence, 0.9)  
smoothed = α * prev + (1 - α) * current
```

- Prevents **flickering** on dashed lanes

3.8 Overlay & HUD (`OverlayRenderer`)

- **Green polyline** = Left detected
- **Blue polyline** = Right detected
- **Gray/dashed** = Low confidence
- **Filled polygon** = Ego-lane corridor
- **HUD** (top-left, large font):

```
Road Lane Assist by: Kipchirchir Raphael, LGL7CS  
Left: YES | Conf: 0.92  
Right: YES | Conf: 0.88  
Lat Offset: -0.05m
```

3.9 Metrics & CSV (`MetricsCalculator`)

- **Lateral offset:** `(vehicle_center - lane_center) * xm_per_pix`
- **CSV columns:**

```
frame_id, left_detected, right_detected, left_conf, right_conf, lat_offset_m
```

4. Results (Real Test Data)

Summary Across 4 Test Videos

Metric	Average	Range
Total Frames	1,392	477 – 1,552
Left Detection Rate	0.492 (49.2%)	20.9% – 61.1%
Right Detection Rate	0.513 (51.3%)	16.8% – 76.7%
Avg Left Confidence	0.507	0.355 – 0.612
Avg Right Confidence	0.552	0.303 – 0.644
Mean Lateral Offset	-0.32 m	-1.06 → +0.89 m
Offset Std Dev	17.8 m	11.4 – 21.6 m
Left Flicker Rate	0.094 / 100 frames	0.042 – 0.136
Right Flicker Rate	0.085 / 100 frames	0.056 – 0.118

Note: Detection rates below 90% due to **challenging conditions**
(shadows, faded paint, curves, glare)

Individual Test Results

Video	Frames	Left Det.	Right Det.	Left Conf.	Right Conf.	Offset (m)	Std Dev
Test 1	477	31.5%	76.7%	0.295	0.644	-0.89	1.14
Test 2	1,083	60.2%	53.3%	0.612	0.572	-2.58	21.6
Test 3	539	61.1%	40.4%	0.474	0.411	+0.75	13.9
Test 4	1,552	20.9%	16.8%	0.355	0.303	-1.06	15.8

5. Qualitative Analysis

Scenario	Observed	Notes
Clear highway	Partial success	Right lane strong, left often missing
Curved road	Unstable	2nd-order poly underfits
Faded paint	Low confidence	Below threshold → "NO"
Shadows	Major failure	Color thresholds fail
Dashed lanes	Flickering	Smoothing helps but not enough
Lane change	Jumps	No prior tracking

6. Failure Modes & Root Causes

Failure	Root Cause	Evidence
Left lane missing	No pixels in histogram	Found 0 left lane pixels
High offset variance	Unstable fits	Std dev up to 21.6 m
Low confidence	Too few pixels / high residual	Pixel: 1.0, Residual: 0.835
Flickering	Confidence near threshold	Final: 0.511, borderline

FICKERINGY CONFIDENCE THRESHOLD: → BORDERLINE

LKA Safety Logic:

If $\text{conf} < 0.4$ for 5+ frames → disengage, alert driver

7. Classical CV vs. Deep Learning

Feature	Classical CV	Deep Learning
Speed	25–40 FPS	10–20 FPS
Robustness	Poor in shadows/faded	Better in all conditions
Interpretability	High	Low
Training	None	Required
Deployment	Embedded	GPU needed

Classical CV failed in real-world variability

DL would likely achieve ≥90% detection

8. Discussion & Recommendations

Key Insights

- Preprocessing is too conservative → filters out valid lane pixels
- ROI + IPM stable, but lane search fails when one side is missing
- Temporal smoothing works, but cannot recover from total loss
- Confidence scoring is accurate but too strict

Recommended Improvements

1. Loosen color thresholds (adaptive per frame)
2. Use 3rd-order polynomial for curves
3. Add fallback lane model (e.g., average lane width)
4. Implement lane tracking (Hungarian + Kalman)
5. Add CLAHE for shadow compensation
6. Try DL model (LaneNet/Ultra-Fast) for comparison

9. Conclusion

The **LKA system** was successfully implemented using **classical CV** and meets **functional requirements**:

- Correct HUD, overlay, CSV
- Real-time processing

- Real-time processing
- Confidence-based detection

However, real-world performance is poor:

- Detection rate: ~50%
- High lateral offset variance
- Frequent flickering

Conclusion:

**Classical CV is insufficient for robust LKA in diverse conditions.
Deep learning or hybrid approach is required for production.**

10. Deliverables (Submitted)

File	Description
<code>outputs/annotated_output.mp4</code>	Full annotated video
<code>outputs/per_frame_metrics.csv</code>	Frame-level metrics
<code>src/</code>	Complete source code
<code>README.md</code>	This report

11. References

1. OpenCV Documentation – `cv2.warpPerspective`, `cv2.polyline`
 2. "Advanced Lane Finding" – Udacity SDC Nanodegree
 3. TuSimple Lane Detection Benchmark
 4. CULane Dataset
-